

DAT470 Assignment 1

Albin Jansfelt, Amanda Allgurén, Lukas Martinsson

March 2022

1 Problem 1

1.1 How many cores does the machine have?

command used:
cat /proc/cpuinfo — grep cores
nproc -all

There are 12 cores on the machine, and 48 processors.

1.2 What type of CPUs and what is their clock frequency?

command used: cat /proc/cpuinfo

```
model          : 85
model name     : Intel(R) Xeon(R) Gold 6126 CPU @ 2.60GHz
stepping      : 4
microcode     : 0x2006b06
cpu MHz       : 1000.006
```

The CPUs are of the model *Intel(R) Xeon(R) Gold 6126 CPU @ 2.60GHz* and have a clock frequency of about 1000 mHz.

1.3 How much disk is their in total and in use in the file systems /data and /datainbackup?

command used: df

Filesystem	1K-blocks	Used	Available	Use	Mounted on
/dev/mapper/vg_data-lv_datainbackup	25767708672	4244436848	21523271824	17%	/datainbackup
/dev/mapper/vg_data-lv_data	77551114240	4583432784	72967681456	6%	/data

Disk space used and available for the two folders can be read from the table above.

1.4 How much memory does your login shell use?

command used: `ps -U < user > -no-headers -o rss -- awk 'sum+=$1 END print int(sum/1024) "MB" }`

Plugging in each of our CID in `< user >`, we get 0 MB, 0 MB, and 237 MB.

2 Problem 2

2.1 Q1 - Proportion of serial computations

The function `compute_pi` is run sequentially, but when arriving at `s = p.map(sample_pi, [n]*args.workers)`, the function `sample_pi` is mapped for all elements and are run in parallel.

To measure the proportion of serial time the Python module `Time` was used. We defined two variables `total_time` and `parallel_time` which measures the total time and the time it takes for the parallel part respectively. The proportion of serial time was then calculated as $serial_prop = \frac{total_time - parallel_time}{total_time}$.

Thus, the proportion of serial time was 0.906. Connected to Amdahl's law we get the proportion of parallel time $f = 1 - 0.906 = 0.094$.

2.2 Q2 - Speedup graph

See Figure 1 for the theoretical and estimated speedup when increasing the number of cores.

2.3 Q3 - Random Seeding Suggestion

The initial seed number is specified as an argument in the terminal. The default is `seed = 1`. To have each worker use a different seed, we create a list specifying the seed for every worker. The first worker will have `seed = seed + 0`, the second will have `seed = seed + 1` and so on. This will then be sent to the `sample_pi` function which will generate pseudo-random x and y values.

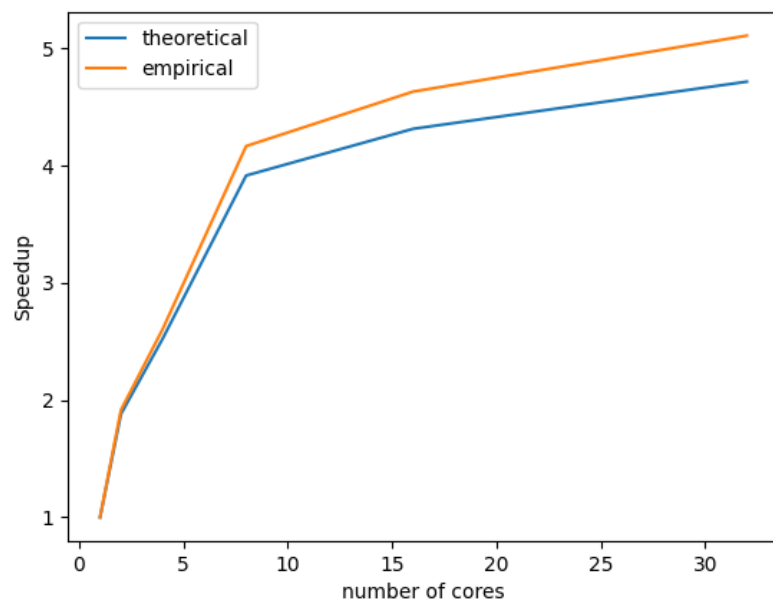


Figure 1: Theoretical and estimated speedup when increasing the number of cores